

OQNLP and MSNLP

Optimal Methods Inc, 7134 Valburn Dr., Austin, TX 78731 www.optimalmethods.com, 512-346-7837

OptTek System, Inc., 1919 7th St., Boulder, CO 80302, www.opttek.com, 303-447-3255

Contents

1	Introduction	337
2	Combining Search Methods and Gradient-Based NLP Solvers	339
3	Output	339
3.1	Log File	339
3.2	The LOCALS File	341
4	The Options File	341
5	Use as a Callable System	345
	Appendix A: Description of the Algorithm	345
	Appendix B: Pure and "Smart" Random Drivers	347
	References	348

1 Introduction

OQNLP and MSNLP are multistart heuristic algorithms designed to find global optima of smooth constrained nonlinear programs (NLPs). By "multistart" we mean that the algorithm calls an NLP solver from multiple starting points, keeps track of all feasible solutions found by that solver, and reports back the best of these as its final solution. The starting points are computed by a scatter search implementation called OptQuest (see www.opttek.com and [Laguna and Marti, 2003]) or by a randomized driver, which generates starting points using probability distributions. There are currently two randomized drivers, Pure Random and Smart Random-see the description of the POINT_GENERATION keyword in Section 3, and Appendix B. With OQNLP, all three drivers are provided, while OptQuest is not present in MSNLP. When interfaced with the GAMS modeling language, any GAMS NLP solver can be called. When used as a callable system, MSNLP uses the LSGRG2 NLP solver (see www.optimalmethods.com and (Smith and Lasdon, 1992)), and this is also provided (optionally) in the GAMS version.

Only the OptQuest driver can handle discrete variables, so OQNLP can attack problems with some or all discrete variables, but MSNLP cannot solve problems with discrete variables. If all variables in a problem are discrete, OQNLP can be applied, but the NLP solver calls play no role, since such solvers vary only the continuous variables, and there aren't any. Thus the OPTQUEST_ONLY option (see details of options in Section 4) should be used. If a problem is nonsmooth (discontinuous functions and/or derivatives, GAMS problem type DNLP), the NLP solver calls may be less reliable than if the problem was smooth. The OPTQUEST_ONLY option may also be useful in this case.

There is no guarantee that the final solution is a global optimum, and no bound is provided on how far that solution is from the global optimum. However, the algorithm has been tested extensively on 135 problems from the set gathered by Chris Floudas [Floudas, et al., 1999], and it found the best known solution on all but three of them to within a percentage gap of 1default parameters and options (which specify 1000 iterations). It solved two of those three to within 1library to within

1seven remaining ones by increasing the iteration limit or using another NLP solver. These results are described in [Lasdon et al., 2004]. For more information on OQNLP, see [Ugray, et. al., 2003].

A multistart algorithm can improve the reliability of any NLP solver, by calling it with many starting points. If you have a problem where you think the current NLP solver is failing to find even a local solution, choose an NLP solver and a limit on the number of solver calls, and try OQNLP or MSNLP. Even if a single call to the solver fails, multiple calls from the widely spaced starting points provided by this algorithm have a much better chance of success.

Often an NLP solver fails when it terminates at an infeasible solution. In this situation, the user is not sure if the problem is really infeasible or if the solver is at fault (if all constraints are linear or convex the problem is most likely infeasible). A multistart algorithm can help in such cases. To use it, the problem can be solved in its original form, and some solver calls may terminate with feasible solutions. The algorithm will return the best of these. If all solver calls terminate infeasible, the problem can be reformulated as a feasibility problem. That is, introduce "deviation" or "elastic" variables into each constraint, which measure the amount by which it is violated, and minimize the sum of these violations, ignoring the true objective. OQNLP or MSNLP can be applied to this problem, and either has a much better chance of finding a feasible solution (if one exists) than does a single call to an NLP solver. If no feasible solution is found, you have much more confidence that the problem is truly infeasible.

The OptQuest and randomized drivers generate trial points which are candidate starting points for the NLP solver. These are filtered to provide a smaller subset from which the solver attempts to find a local optimum. In the discussion which follows, we refer to this NLP solver as "L.", for Local solver.

The most general problem OQNLP can solve has the form

$$\text{minimize } f(x,y) \quad (1.1)$$

subject to the nonlinear constraints

$$gl \leq G(x,y) \leq gu \quad (1.2)$$

and the linear constraints

$$l \leq A_1x + A_2y \leq u \quad (1.3)$$

$$x \in S, \quad y \in Y \quad (1.4)$$

where x is an n -dimensional vector of continuous decision variables, y is a p -dimensional vector of discrete decision variables, and the vectors gl, gu, l , and u contain upper and lower bounds for the nonlinear and linear constraints respectively. The matrices A_1 and A_2 are m_2 by n and m_2 by p respectively, and contain the coefficients of any linear constraints. The set S is defined by simple bounds on x , and we assume that it is closed and bounded, i.e., that each component of x has a finite upper and lower bound. This is required by all drivers (see section 4 for a discussion of the parameter ARTIFICIAL_BOUND which provides bounds when none are specified in the model). The set Y is assumed to be finite, and is often the set of all p -dimensional binary or integer vectors y . The objective function f and the m_1 - dimensional vector of constraint functions G are assumed to have continuous first partial derivatives at all points in $S \times Y$. This is necessary so that L can be applied to the relaxed NLP sub-problems formed from 1.1 - 1.4 by allowing the y variables to be continuous. The MSNLP system does not allow any discrete variables.

An important function used in this multistart algorithm is the L_1 exact penalty function, defined as

$$P_1(x, w) = f(x) + \sum_{i=1}^m w_i \text{viol}(g_i(x)) \quad (1.5)$$

where the w_i are nonnegative penalty weights, $m = m_1 + m_2$, and the vector g has been extended to include the linear constraints 1.4. For simplicity, we assume there are no y variables: these would be fixed when this function is used. The function $\text{viol}(g_i(x))$ is equal to the absolute amount by which the i th constraint is violated at the point x . It is well known (see [Nash and Sofer, 1996]) that if x^* is a local optimum of 1.1 - 1.4, u^* is a corresponding optimal multiplier vector, the second order sufficiency conditions are satisfied at (x^*, u^*) , and

$$w_i > \text{abs}(u_i^*) \quad (1.6)$$

then x^* is a local unconstrained minimum of P_1 . If 1.1 - 1.4 has several local minima, and each w_i

is larger than the maximum of all absolute multipliers for constraint i over all these optima, then P_i has a local minimum at each of these local constrained minima. We will use P_i to set thresholds in the merit filter.

2 Combining Search Methods and Gradient-Based NLP Solvers

For smooth problems, the relative advantages of a search method over a gradient-based NLP solver are its ability to locate an approximation to a good local solution (often the global optimum), and the fact that it can handle discrete variables. Gradient-based NLP solvers converge to the "nearest" local solution, and have no facilities for discrete variables, unless they are imbedded in a rounding heuristic or branch-and-bound method. Relative disadvantages of search methods are their limited accuracy, and their weak abilities to deal with equality constraints (more generally, narrow feasible regions). They find it difficult to satisfy many nonlinear constraints to high accuracy, but this is a strength of gradient-based NLP solvers. Search methods also require an excessive number of iterations to find approximations to local or global optima accurate to more than two or three significant figures, while gradient-based solvers usually achieve four to eight-digit accuracy rapidly. The motivation for combining search and gradient-based solvers in a multi-start procedure is to achieve the advantages of both while avoiding the disadvantages of either.

3 Output

3.1 Log File

When it operates as a GAMS solver, OQNLP and MSNLP will by default write information on their progress to the GAMS log file. When used as a callable system, this information, if requested, will be written to a file opened in the users calling program. The information written consists of:

1. Echos of important configuration and setup values
2. Echo (optionally) of options file settings processed
3. Echos of important algorithm settings, parameters, and termination criteria
4. The iteration log
5. Final results, termination messages, and status report

A segment of that iteration log from stages 1 and 2 of the algorithm is shown below for the problem *ex8.6.2_30.gms*, which is one of a large set of problems described in [Floudas, et al., 1999]. This is a 91 variable unconstrained minimization problem, available from GLOBALLib at www.gamsworld.org/global. There are 200 iterations in stage one and 1000 total iterations (see Appendix A for an algorithm description), with output every 20 iterations and every solver call.

The headings below have the following meanings:

Itn	iteration number
Penval	Penalty function value
Merit Filter	ACC if the merit filter accepts the point, REJ if it rejects
Merit	threshold value for merit filter: accepts if Penval < Threshold
Threshold	
Dist Filter	ACC if the distance filter accepts the point, REJ if it rejects
Best Obj	Best feasible objective value found thus far
Solver Obj	Objective value found by NLP solver at this iteration
Term Code	Code indicating reason for termination of NLP solver: KTC means Kuhn-Tucker optimality conditions satisfied FRC means that the fractional objective change is less than a tolerance for some number of consecutive iterations INF means solver stopped at an infeasible point
Sinf	sum of infeasibilities at point found by NLP solver

Iterations 0 through 200 below show the initial NLP solver call (at the user-specified initial point, which finds a local minimum with objective value -161.8), and every 20th iteration of stage 1, which has no other solver calls. At iteration 200 stage 1 ends, and the solver is started at the best of the 200 stage 1 points, finding a local min with objective -176.0. The next solver call at iteration 207 finds a better objective of -176.4. Note that, at iteration 207, the OptQuest trial solution has a Penval of -23.18, and this is less than the merit threshold of -20.75, so the merit filter ACCEpts the trial solution, as does the distance filter. The next 9 solver calls fail to improve this value, so Best Obj remains the same, until at iteration 432 a solution with value -176.6 is found. At iteration 473, the solver call finds a value of -177.5. Further solver calls do not find an improved solution and are not shown. The solution with value -177.5 is the best known solution, but OQNLP cannot guarantee this.

Itn	Penval	Merit	Merit	Dist	Best	Solver	Term	Sinf
		Filter	Threshold	Filter	Obj	Obj	Code	
0	+1.000e+030		-1.000e+030		-1.618e+002	-1.618e+002	FRC	+0.000e+000
20	-4.485e+000							
40	-6.321e+000							
60	-1.126e+001							
80	+2.454e+000							
100	+8.097e+001							
120	+5.587e+001							
140	+1.707e+004							
160	+2.034e+002							
180	+7.754e+001							
200	-6.224e+000							

Itn	Penval	Merit	Merit	Dist	Best	Solver	Term	Sinf
		Filter	Threshold	Filter	Obj	Obj	Code	
201	+1.000e+030	ACC	-1.000e+030	ACC	-1.618e+002	-1.760e+002	FRC	+0.000e+000
207	-2.318e+001	ACC	-2.075e+001	ACC	-1.760e+002	-1.764e+002	FRC	+0.000e+000
220	-8.324e+000	REJ	-2.318e+001	ACC	-1.764e+002			
240	+8.351e+000	REJ	-1.834e+001	ACC	-1.764e+002			
251	-1.117e+001	ACC	-1.008e+001	ACC	-1.764e+002	-1.682e+002	FRC	+0.000e+000
256	-1.244e+001	ACC	-1.117e+001	ACC	-1.764e+002	-1.758e+002	FRC	+0.000e+000
258	-1.550e+001	ACC	-1.244e+001	ACC	-1.764e+002	-1.678e+002	FRC	+0.000e+000
260	-7.255e+000	REJ	-1.550e+001	ACC	-1.764e+002			
280	+8.170e+001	REJ	-1.220e+001	ACC	-1.764e+002			
282	-2.521e+001	ACC	-1.220e+001	ACC	-1.764e+002	-1.758e+002	FRC	+0.000e+000
300	+5.206e+001	REJ	-2.521e+001	ACC	-1.764e+002			
300	+5.206e+001	REJ	-2.521e+001	ACC	-1.764e+002			
320	+1.152e+000	REJ	-1.642e+001	ACC	-1.764e+002			
329	-2.111e+001	ACC	-1.294e+001	ACC	-1.764e+002	-1.763e+002	FRC	+0.000e+000
338	-3.749e+001	ACC	-2.111e+001	ACC	-1.764e+002	-1.763e+002	FRC	+0.000e+000
340	+2.235e+002	REJ	-3.749e+001	ACC	-1.764e+002			
360	+8.947e+001	REJ	-2.363e+001	ACC	-1.764e+002			
366	-3.742e+001	ACC	-2.363e+001	ACC	-1.764e+002	-1.761e+002	FRC	+0.000e+000
380	-2.244e+001	REJ	-3.742e+001	ACC	-1.764e+002			
391	-2.974e+001	ACC	-2.244e+001	ACC	-1.764e+002	-1.754e+002	FRC	+0.000e+000
400	+1.986e+002	REJ	-2.974e+001	ACC	-1.764e+002			
400	+1.986e+002	REJ	-2.974e+001	ACC	-1.764e+002			
420	-1.231e+001	REJ	-2.359e+001	ACC	-1.764e+002			
432	-2.365e+001	ACC	-2.359e+001	ACC	-1.764e+002	-1.766e+002	FRC	+0.000e+000
440	+6.335e+000	REJ	-2.365e+001	ACC	-1.766e+002			
460	-8.939e+000	REJ	-1.872e+001	ACC	-1.766e+002			
473	-3.216e+001	ACC	-1.872e+001	ACC	-1.766e+002	-1.775e+002	FRC	+0.000e+000
480	+1.744e+002	REJ	-3.216e+001	ACC	-1.775e+002			

3.2 The LOCALS File

The LOCALS file is a text file containing objective and variable values for all local solutions found by MSNLP. It is controlled by the LOCALS_FILE and LOCALS_FILE_FORMAT keywords in the MSNLP Options file. An example for the problem EX_8.1_5 from the Floudas problem set (available on www.gamsworld.org, link to globalworld) is shown below. The headings, included for explanatory purposes and not part of the file, have the following meaning:

No. index of local solution

Obj objective value of local solution

Var variable index

Value variable value

No.	Obj	Var	Value
1	-1.03163e+000	1	-8.98448e-002
1	-1.03163e+000	2	7.12656e-001
2	-1.03163e+000	1	8.98418e-002
2	-1.03163e+000	2	-7.12656e-001
3	-2.15464e-001	1	1.70361e+000
3	-2.15464e-001	2	-7.96084e-001
4	-2.15464e-001	1	-1.70361e+000
4	-2.15464e-001	2	7.96084e-001
5	0.00000e+000	1	0.00000e+000
5	0.00000e+000	2	0.00000e+000
6	2.10425e+000	1	1.60710e+000
6	2.10425e+000	2	5.68656e-001
7	2.10425e+000	1	-1.60711e+000
7	2.10425e+000	2	-5.68651e-001

Thus local solutions 1 and 2 both have objective values of -1.03163. The first solution has variable values $x = -8.98448e-002$, $y = 7.12656e-001$, where these are in the same order as they are defined in the gams model. The second local solution has $x = 8.98418e-002$, $y = -7.12656e-001$. Seven local solutions are found. This output is produced with all default parameter values for MSNLP options and tolerances, except the distance and merit filters were turned off, i.e the keywords USE_DISTANCE_FILTER and USE_MERIT_FILTER were set to 0 in the MSNLP options file. This causes the NLP solver to be called at every stage 2 trial point, and is recommended if you wish to obtain as many local solutions as possible.

4 The Options File

The options file is a text file containing a set of records, one per line. Each record has the form <keyword> <value>, where the keyword and value are separated by one or more spaces. All relevant options are listed in this guide. You can also get a sample option file with all options and their default values by specifying the single option help in an option file. The list of all options appears in the log file. The options are described below.

Option	Description	Default
ARTIFICIAL_BOUND	This value (its negative) is given to the driver as the upper (lower) bound for any variable with no upper or lower bound. However, the original bounds are given to the local solver, so it can produce solutions not limited by this artificial bound. All drivers must have finite upper and lower bounds for each variable. If ARTIFICIAL_BOUND (or any of the user-supplied bounds) is much larger than any component of the optimal solution, the driver will be less efficient because it is searching over a region that is much larger than needed. Hence the user is advised to try to provide realistic values for all upper and lower bounds. It is even more dangerous to make ARTIFICIAL_BOUND smaller than some component of a globally optimal solution, since the driver can never generate a trial point near that solution. It is possible, however, for the local solver to reach a global solution in this case, since the artificial bounds are not imposed on it.	1.e4
BASIN_DECREASE_FACTOR	This value must be between 0 and 1. If DYNAMIC_DISTANCE_FILTER is set to 1, the MAXDIST value associated with any local solution is reduced by (1-BASIN_DECREASE_FACTOR) if WAITCYCLE consecutive trial points have distance from that solution less than MAXDIST.	0.2
BASIN_OVERLAP_FIX	A value of 1 turns on logic which checks the MAXDIST values of all pairs of local solutions, and reduces any pair of MAXDIST values if their sum is greater than the distance between the 2 solutions. This ensures that the spherical models of their basins of attracting do not overlap. A value of 0 turns off this logic. Turning it off can reduce the number of NLP solver calls, but can also cause the algorithm to miss the global solution.	1
DISTANCE_FACTOR	If the distance between a trial point and any local solution found previously is less than DISTANCE_FACTOR * MAXDIST, the NLP solver is not started from that trial point. MAXDIST is the largest distance ever traveled to get to that local solution. Increasing DISTANCE_FACTOR leads to fewer solver calls and risks finding a worse solution. Decreasing it leads to more solver calls and possibly a better solution.	1.0
DYNAMIC_DISTANCE_FILTER	A value of 1 turns on logic which reduces the value of MAXDIST (described under the USE_DISTANCE_FILTER keyword) for a local solution if WAITCYCLE consecutive trial points have a their distances from that solution less than MAXDIST. MAXDIST is multiplied by (1-BASIN_REDUCTION_FACTOR). A value of 0 turns off this logic. Turning it off can decrease the number of NLP solver calls, but can also lead to a worse final solution.	1
DYNAMIC_MERIT_FILTER	A value of 1 turns on logic which dynamically varies the parameter which increases the merit filter threshold, THRESHOLD_INCREASE_FACTOR. If WAITCYCLE consecutive trial points have been rejected by the merit filter, this value is replaced by max(THRESHOLD_INCREASE_FACTOR, val), where val is the value of THRESHOLD_INCREASE_FACTOR which causes the merit filter to just accept the best of the previous WAITCYCLE trial points. A value of 0 turns off this logic. Turning it off can reduce NLP solver calls, but may lead to a worse final solution.	1
ENABLE_SCREEN_OUTPUT	A value of 0 turns off the writing of the iteration log and termination messages to the gams log file that appears on the screen, while 1 enables it.	1
ENABLE_STATISTICS_LOG	Using a value of 1 creates a text file called stats.log in the project directory containing one line of problem (name, variables, constraints) and performance information (best objective value, total solver time, iterations, iterations to best solution, etc) for each problem solved.	0

Option	Description	Default
FEASIBILITY TOLERANCE	This tolerance is used to check each point returned by an NLP solver for feasibility. If the largest absolute infeasibility at the point is larger than this tolerance, the point is classified infeasible. This test is made because points returned by NLP solvers may occasionally be infeasible despite feasible status codes. Some NLP solvers use internal scaling before testing for feasibility. The unscaled problem may be infeasible, while the scaled one is feasible. If this occurs, increasing this tolerance (to 1.e-2 or larger) often eliminates the problem.	1.e-4
ITERATION_PRINT_ FREQUENCY	If the OQNLP iteration log is written to the GAMS log file, one line of output is written every k 'th OptQuest iteration, where k is the value given here.	20
ITERATION_LIMIT	Increasing this limit can allow OQNLP to find a better solution. Try it if your run using 1000 iterations doesn't take too long. Surprisingly, the best solution using, say 2000 iterations, may be found in the first 1000 iterations, and that solution may be better than the one found with an iteration limit of 1000. This is because OptQuest changes its search strategy depending on the iteration limit. Because of this, it is also possible that increasing the iteration limit will yield a worse solution, but this is rare. Decreasing this iteration limit usually leads to a worse solution, but also reduces run time. OQNLP iterations can not be set using GAMS <code>iterlim</code> . The GAMS <code>iterlim</code> is used as the iteration limit for the NLP subsolves in an OQNLP run	1000
LOCALS_FILE	Specify a complete path and name for a file to which the objective value and values of all variables for all local solutions found will be written. For example, C:\mydirectory\locals.out. There are 2 possible formats for this file, specified by the LOCALS_FILE_FORMAT option below. If there is no LOCALS_FILE record in the options file, the locals file will not be created.	No locals file created
LOCALS_FILE_FORMAT	There are 2 possible values for this option. The REPORT entry creates the locals file in a format designed to be examined easily by eye, but processed less easily by a computer program or spreadsheet. The DATA1 entry creates a file with many records, each on a single line, each line having the following format: <index of local optimum> <objval> <var index> <var value>	REPORT
LOGFILE_ITN_PRINT_ FREQUENCY	In the output written to the log file, one line of output is written every k 'th iteration, where k is the value given here.	20
MAX_LOCALS	When the number of distinct local solutions found exceeds the value specified here, the system will stop, returning the best solution found.	1000
MAX_SOLVER_CALLS	When the number of calls to the NLP solver exceeds the value specified here, the system will stop, returning the best solution found.	1000
MAX_SOLVER_CALLS_ NOIMPROVEMENT	The positive integer specified here will cause the system to stop whenever the number of consecutive solver calls with a fractional improvement in the best objective value found less than 1.e-4 exceeds that value. In other words, if the value specified is 50, and there are more than 50 consecutive solver calls where the relative change in the best objective was less than 1.e-4 in all iterations, the system will stop.	100
MAXTIME	When the execution time exceeds this value, the system will stop, returning the best solution found.	1000 seconds
NLPSOLVER	This option is available only within GAMS. It specifies the NLP solver to be called. Any GAMS NLP solver for which the user has a license can be used. Further, one can specify an option file for the GAMS NLP solver by appending a ".n" with $n=1..999$ to the solver name. For example, NLPSOLVER conopt.1 will instruct the NLP solver CONOPT to use option file conopt.opt, NLPSOLVER conopt.2 will make CONOPT read option file conopt.op2 and so on.	LSGRG

Option	Description	Default
OPTQUEST_ONLY	This option applies only to the OptQuest driver. If you think the NLP solver is taking too long and/or not working well, choosing 1 will stop it from being called. This may occur if the problem is of type "DNLP", where one or more problem functions are discontinuous or have discontinuous derivatives. If the problem has only discrete (integer) variables, choose 1, as there is nothing for the NLP solver to do (since it optimizes over the continuous variables when the integers are fixed, and there aren't any).	0
OQNLP_DEBUG	Values of 1 or 2 cause more information to be written to the iteration log. The default value of 0 suppresses all this output.	0
POINT_GENERATION	OPTQUEST causes trial points to be generated by the OptQuest driver RANDOM causes trial points to be generated by sampling each variable from a uniform distribution defined within its bounds SMARTRANDOM1 generates trial points by sampling each variable independently from either normal or triangular distributions, whose parameters are determined as described in Appendix A.	SMART- RANDOM1 (MSNLP) OPTQUEST (OQNLP)
SAMPLING_ DISTRIBUTION	This keyword is relevant only when POINT_GENERATION is set to SMARTRANDOM1. Then a value of 0 causes normal distributions to be used to generate trial points, while a value of 1 causes triangular distributions to be used.	0
SEARCH_TYPE	This option applies only to the OptQuest driver, and controls the search strategy used by OptQuest. The three choices that are relevant for use within OQNLP are: aggressive This choice controls the population update in step 7 of the OptQuest algorithm (see Appendix A). It triggers a very aggressive update, which keeps the best of the points generated from the current population as the new population. The risk in this is that all points in the new population may cluster in a small portion of the search volume, and regions far from this volume will not be explored in the next cycle. boundary This option affects the trial points generated by OptQuest, directing them toward the boundary of the region defined by the linear constraints and variable bounds. The value of SEARCH_PARAMETER discussed below controls the fraction of points that are directed toward the boundary. crossover This option affects how OptQuest trial points are generated from population points. It retains the linear combination operator discussed in Appendix A, but adds a "crossover" operator, similar to those used in evolutionary or genetic algorithms, to create 2 additional trial points.	boundary
SOLVELINK	This option defines the solvelink used for the NLP solver: 1: Call NLP solver via script 2: Call NLP and MIP solver via module 5: Call NLP and MIP solver in memory	5
SOLVER_LOG_TO_GAMS_ LOG	Setting the parameter to 1 instructs OQNLP to copy the log from the NLP subsolver to the OQNLP log. It can be very helpful to inspect the NLP subsolver log especially if the solver termination code is "???".	0
STAGE1_ITERATIONS	Specifies the total number of iterations in stage 1 of the algorithm, where no NLP solver calls are made. Increasing this sometimes leads to a better starting point for the first local solver call in stage 2, at the cost of delaying that call. Decreasing it can lead to more solver calls, but the first call occurs sooner.	200
THRESHOLD_INCREASE_ FACTOR	This value must be nonnegative. If there are WAITCYCLE (see below) consecutive OptQuest iterations where the merit filter logic causes the NLP solver not to be called, the merit threshold is increased by multiplying it by (1+THRESHOLD_INCREASE_FACTOR)	0.2

Option	Description	Default
USE_DISTANCE_FILTER	Use 0 to turn off the distance filter, the logic which starts the NLP solver at a trial point only if the (Euclidean) distance from that point to any local solution found thus far is greater than the distance threshold. Turning off the distance filter leads to more solver calls and more run time, and increases the chances of finding a global solution. Turn off both distance and merit filters to find (almost) all local solutions.	1
USE_LINEAR_CONSTRAINTS	This option applies only to the OptQuest driver, and to problems that have linear constraints other than simple bounds on the variables. Using 1 (all OptQuest trial points satisfy the linear constraints) often leads to fewer iterations and solver calls, but OptQuest has to solve an LP to project each trial point onto the linear constraints. For large problems (more than 100 variables), this can greatly increase run time, so the default value is off (0).	0
USE_MERIT_FILTER	Use 0 to turn off the merit filter, the logic which starts the NLP solver at a trial point only if the penalty function value at that point is below the merit threshold. This will lead to more solver calls, but increases the chances of finding a global solution. Turn off both filters if you want to find (almost) all local solutions. This will cause the solver to be called at each stage 2 iteration.	1
WAITCYCLE	This value must be a positive integer. If the merit filter is used, and there are WAITCYCLE consecutive iterations where the merit filter logic causes the NLP solver not to be started, the merit filter threshold is increased by the factor THRESHOLD_INCREASE_FACTOR (see above). Increasing WAITCYCLE usually leads to fewer solver calls, but risks finding a worse solution. Decreasing it leads to more solver calls, but may find a better solution.	20

5 Use as a Callable System

MSNLP and OQNLP is also available as a callable system. It currently uses the LSGRG2 NLP solver as its local solver, but any other NLP solver could be included. A sample calling program is provided which a user can easily adapt. The user must provide a C function which computes values of the objective and all constraint functions, given current values of all variables. First partial derivatives of these functions can be approximated by forward or central differences, or may be computed in a user-provided function.

Appendix A: Description of the Algorithm

A pseudo-code description of the MSNLP algorithm follows, in which $SP(xt)$ denotes the starting point generator and xt is the candidate starting point produced. We refer to the local NLP solver as $L(xs, xf)$, where xs is the starting point and xf the final point. The function $UPDATE_LOCALS(xs, xf, w)$ processes and stores solver output xf , using the starting point xs to compute the distance from xs to xf , and produces updated penalty weights, w . For more details, see [Lasdon, Plummer et al., 2004].

MSNLP Algorithm

STAGE 1

x_0 = user initial point

Call $L(x_0, xf)$

Call $UPDATE_LOCALS(x_0, xf, w)$

FOR $i = 1, n1$ **DO**

 Call $SP(xt(i))$

 Evaluate $P(xt(i), w)$

ENDDO

xt^* = point yielding best value of $P(xt(i), w)$ over all stage one points, $(i = 1, 2, \dots, n1)$.

call $L(xt^*, xf)$

Call UPDATE LOCALS(xt^*, xf, w)

threshold = $P(xt^*, w)$

STAGE 2

FOR $i = 1, n2$ **DO**

 Call SP($xt(i)$)

 Evaluate $P(xt(i), w)$

 Perform merit and distance filter tests:

 Call distance filter($xt(i)$, dstatus)

 Call merit filter($xt(i)$, threshold, mstatus)

IF (dstatus and mstatus = "accept") **THEN**

 Call $L(xt(i), xf)$

 Call UPDATE LOCALS($xt(i), xf, w$)

ENDIF

ENDDO

After an initial call to L at the user-provided initial point, x_0 , stage 1 of the algorithm performs $n1$ iterations in which SP(xt) is called, and the L1 exact penalty value $P(xt, w)$ is calculated. The user can set $n1$ through the MSNLP options file using the STAGE1_ITERATIONS keyword. The point with the smallest of these P values is chosen as the starting point for the next call to L , which begins stage 2. In this stage, $n2$ iterations are performed in which candidate starting points are generated and L is started at any one which passes the distance and merit filter tests. The options file keyword STAGE2_ITERATIONS sets $n2$.

The distance filter helps insure that the starting points for L are diverse, in the sense that they are not too close to any previously found local solution. Its goal is to prevent L from starting more than once within the basin of attraction of any local optimum. When a local solution is found, it is stored in a linked list, ordered by its objective value, as is the Euclidean distance between it and the starting point that led to it. If a local solution is located more than once, the maximum of these distances, $maxdist$, is updated and stored. For each trial point, t , if the distance between t and any local solution already found is less than $DISTANCE_FACTOR * maxdist$, L is not started from the point, and we obtain the next trial solution from the generator.

This distance filter implicitly assumes that the attraction basins are spherical, with radii at least $maxdist$. The default value of $DISTANCE_FACTOR$ is 1.0, and it can be set to any positive value in the MSNLP options file-see Section 3. As $DISTANCE_FACTOR$ approaches zero, the filtering effect vanishes, as would be appropriate if there were many closely spaced local solutions. As it becomes larger than 1, the filtering effect increases until eventually L is never started.

The merit filter helps insure that the starting points for L have high quality, by not starting from candidate points whose exact penalty function value P_1 (see equation (5), Section 1) is greater than a threshold. This threshold is set initially to the P_1 value of the best candidate point found in the first stage of the algorithm. If trial points are rejected by this test for more than WAITCYCLE consecutive iterations, the threshold is increased by the updating rule:

$$\text{threshold} \leftarrow \text{threshold} + \text{THRESHOLD_INCREASE_FACTOR} * (1.0 + \text{abs}(\text{threshold}))$$

where the default value of $\text{THRESHOLD_INCREASE_FACTOR}$ is 0.2 and that for WAITCYCLE is 20. The additive 1.0 term is included so that threshold increases by at least $\text{THRESHOLD_INCREASE_FACTOR}$ when its current value is near zero. When a trial point is accepted by the merit filter, threshold is decreased by setting it to the P_1 value of that point.

The combined effect of these 2 filters is that L is started at only a few percent of the trial points, yet global optimal solutions are found for a very high percentage of the test problems. However, the chances of finding a global optimum are increased by increasing ITERATION_LIMIT (which we recommend trying first) or by "loosening" either or both filters, although this is rarely necessary in our tests if the dynamic filters and basin overlap fix are used, as they are by default. If the ratio of stage 2 iterations to solver calls is more than 20 using the current filter parameters, and computation times with the default filter parameters are reasonable, you can try loosening the filters. This is achieved for the merit filter either by decreasing

WAITCYCLE or by increasing THRESHOLD_INCREASE_FACTOR (or doing both), and for the distance filter by decreasing DISTANCE_FACTOR. Either or both filters may be turned off, by setting USE_DISTANCE_FILTER and/or USE_MERIT_FILTER to 0. Turning off both causes an NLP solver call at every stage 2 trial point. This is the best way to insure that all local optima are found, but it can take a long time.

Appendix B: Pure and "Smart" Random Drivers

The "pure" random (PR) driver generates uniformly distributed points within the hyper-rectangle S defined by the variable bounds. However, this rectangle is often very large, because users often set bounds to $(-\infty, +\infty)$, $(0, +\infty)$, or to large positive and/or negative numbers, particularly in problems with many variables. This usually has little adverse impact on a good local solver, as long as the starting point is chosen well inside the bounds. But the PR generator will often generate starting points with very large absolute component values when some bounds are very large, and this sharply degrades solver performance. Thus we were motivated to develop random generators which control the likelihood of generating candidate points with large components, and intensify the search by focusing points into promising regions. We present two variants, one using normal, the other triangular distributions. Pseudo-code for this "smart random" generator using normal distributions follows, where w is the set of penalty weights determined by the "update locals" logic discussed above, after the first solver call at the user-specified initial point.

Smart Random Generator with Normal Distributions, $SRN(xt)$

IF (first call) **THEN**

Generate $k1$ (default 400) diverse points in S and evaluate the exact penalty function $P(x, w)$ at each point.

B = subset of S with $k2$ (default 10) best P values

FOR $i = 1, nvars$ **DO**

$xmax(i) = \max$ of component i of points in B

$xmin(i) = \min$ of component i of points in B

$\mu(i) = (xmax(i) + xmin(i))/2$

$ratio(i) = (xmax(i) - xmin(i))/(1 + buvar(i) - blvar(i))$

$sigfactor = 2.0$

IF ($ratio > 0.7$) $sigfactor = f(ratio)$

$\sigma(i) = (xmax(i) - xmin(i))/sigfactor$

ENDDO

ENDIF

FOR $i = 1, nvars$ **DO**

Generate a normally distributed random variable $rv(i)$ with mean $\mu(i)$ and standard deviation $\sigma(i)$

If $rv(i)$ is between $blvar(i)$ and $buvar(i)$, $xt(i) = rv(i)$

If $rv(i) < blvar(i)$, generate $xt(i)$ uniformly between $blvar(i)$ and $xmin(i)$

If $rv(i) > buvar(i)$, generate $xt(i)$ uniformly between $xmax(i)$ and $buvar(i)$

ENDDO

Return xt

This SRN generator attempts to find a subset, B , of $k2$ "good" points, and generates most of its trial points xt , within the smallest rectangle containing B . It first generates a set of $k1$ diverse points within the bounds using a stratified random sampling procedure with frequency-based memory. For each variable $x(i)$, this divides the interval $[blvar(i), buvar(i)]$ into 4 equal segments, chooses a segment with probability inversely proportional to the frequency with which it has been chosen thus far, then generates a random point in this segment. We choose $k2$ of these points having the best $P(x, w)$ penalty values, and use the smallest rectangle containing these, intersecting the i th axis at points $[xmin(i), xmax(i)]$, to

define n univariate normal distributions (driver SRN) or n univariate triangular distributions (driver SRT). The mean of the i th normal distribution, $\mu(i)$, is the midpoint of the interval $[x_{\min}(i), x_{\max}(i)]$, and this point is also the mode of the i th triangular distribution, whose lower and upper limits are $blvar(i)$ and $buvar(i)$. The standard deviation of the i th normal distribution is selected as described below. The trial point xt is generated by sampling n times independently from these distributions. For the driver using normals, if the generated point lies within the bounds, it is accepted. Otherwise, we generate a uniformly distributed point between the violated bound and the start of the interval.

To determine the standard deviation of the normal distributions, we compute *ratio*, roughly the ratio of interval width to distance between bounds, where the factor 1.0 is included to avoid division by zero when the bounds are equal (fixed variables). If the interval width is small relative to the distance between bounds for variable i ($ratio \leq 0.7$), then the standard deviation $\sigma(i)$ is half the interval width, so about 1/3 of the $xt(i)$ values fall outside the interval, providing diversity when the interval does not contain an optimal value for $x(i)$. If the bounds are large, then *ratio* should be small, say less than 0.1, so $xt(i)$ values near the bounds are very unlikely. If $ratio > 0.7$, the function f sets sigfactor equal to 2.56 if *ratio* is between 0.7 and 0.8, increasing in steps to 6.2 if $textit{ratio} > 0.999$. Thus if *ratio* is near 1.0, more than 99% of the values fall within the interval, and few have to be projected back within the bounds. The projecting back process avoids undesirable clustering of trial points at a bound, by generating points uniformly between the violated bound and the nearest edge of the interval $[x_{\min}(i), x_{\max}(i)]$. When the interval $[x_{\min}(i), x_{\max}(i)]$ is sharply skewed toward one of the variable bounds and is much narrower than the distance between the bounds, a symmetric distribution like the normal, combined with our projection procedure, generates too many points between the interval and its nearest bound. A quick scan of the test results indicates that this happens rarely, but an asymmetric distribution like the triangular overcomes this difficulty, and needs no projection.

References

- Floudas, C.A., et al. 1999. Handbook of Test Problems in Local and Global Optimization. Kluwer Academic Publishers.
- Laguna, Manuel, R. Marti. 2003. Scatter Search, Methodology and Implementations in C. Kluwer Academic Publishers.
- Lasdon, L., J. Plummer, Z. Ugray, and M. Bussieck. 2004. Improved Filters and Randomized Drivers for Multi-start Global Optimization, working paper, MSIS Department, McCombs College of Business, May 2004. Submitted to Mathematical Programming.
- Nash, S. G., A. Sofer. 1996. Linear and Nonlinear Programming. McGraw-Hill Companies, Inc.
- Smith, S., L. Lasdon. 1992. Solving Large Sparse Nonlinear Programs Using GRG. ORSA Journal on Computing 4 1 3-15.
- Ugray, Z., L. Lasdon, J. Plummer, et.al. 2003. A Multistart Scatter Search Heuristic for Smooth NLP and MINLP problems, submitted to Informs Journal on Computing. Copies available on request from the author, or on the web at www.utexas.edu/courses/lasdon, link to papers.